

Hierarchical goal abstraction for sensorimotor agency

The model – Draft 0.1 May 18, 2016




Figure 1: Architecture of the model.

1 Simulator

- **Environment:** the simulation is done in a 2-dimensional space without physics.
- **Agent:**
 - The whole body is composed of the 2 arms plus a segment joining their origins
 - Each arm has 3 joints (3DoF)
 - 22 touch sensors all over the body.
- **Agent sensory information:**

- **vision (V)**: a 20x20pixels retina on which the body (black line on figure 1) is depicted in the current position.
- **proprioception (P)**: a 20x20pixels retina on which the current angles of joints are represented.

Encoding:

- * Angles are encoded as 2D gaussians on a horizontal line centered in the middle of the retina.
 - * The position of each gaussian on the horizontal line is related to the position of the related joint in the one-dimentional space of the body.
 - * The variance of each gaussian on the y-axis is related to the amplitude of the represented joint angle.
 - * The variance of each gaussian on the x-axis is a proportion of the one on the y-axis, so that overlapping between gaussians is not too much.
- **touch (T)**: a 20x20pixels retina on which the current activations of the touch sensors are represented.

Encoding:

- * Sensors are encoded as 2D gaussians on a horizontal line centered in the middle of the retina.
- * The position of each gaussian on the horizontal line is related to the position of the sensor in the one-dimentional space of the body.
- * The variance of each gaussian on the y-axis is related to the amplitude of the activation of the represented sensor.
- * The variance of each gaussian on the x-axis is a proportion of the one on the y-axis, so that gaussians do not overlap too much.

2 Controller

The controller takes the sensory information at each timestep and outputs a motor command consisting in the current requested position of the 6 joint angles.

2.1 Definitions

- goal:**
- A state of the world that the **goal-abstraction layer** can identify and distinguish from others (section 2.1).
 - A configuration of the **action-selection layer** that can be linked to a proper action (section 2.2).

Note: the goal-abstraction layer has the same dimensionality of

Note: the goal-abstraction layer has the same dimensionality of the goal-selection layer so that they can be compared.

- match:**
- The occurrence of an identical configuration of the goal-abstraction and goal-selection layers.
- prediction:**
- Computation of the probability that given a configuration of the goal-selection layer the agent will obtain a state in which the goal-abstraction layer has the same configuration (*match* event).
- error:**
- Difference between the prediction and the real occurrence of a *match* event within a trial.

The sensory information is further processed by computing the finite difference with the information given at the previous timestep. As a result the actual input is composed by three retinas containing the *derivatives* of the (V), (P) and (T) retinas described above.

The result of this processing is sent as input to a 4-layered hierarchical set of self-organizing maps (SOMs). The output of each SOM is further processed so that all output units have 0 activation except the one whose weights are the closest to the current input, which has 1 activation. There is a further threshold based on the current learning within the SOM, so that if the weights for the cluster to which the current input should belong have not been learned enough then all output units have 0 activation.

In the 1st layer of the hierarchical network the segregation between sensory modalities is maintained. Thus this layer is composed of three SOMs each reducing the 20x20pixel input of a sensory modality to a 8x8pixel output.




Figure 2: Activity of the 1st layer of the hierarchical network. A) The current position of the agent. B) The first row shows the derivatives of the (V),(P) and (T) retinas. The second row show the current status of the weights of the SOMs. Each of the three plots shows a 8x8 matrix of 20x20 retinas. Each 20x20 retina represents the weights connecting the input retina to an output unit. The third row shows the output layers of each SOM. the black pixel represents the winner unit for the current input.

The 2nd layer of the hierarchical network is composed of two SOMs. The 1st SOM in the 2nd layer merges the 16x8pixel output coming from the (V) and the (P) SOMs in the first layer into a 4x4pixel output. The 2st SOM in the 2nd layer merges the 16x8pixel output coming from the (P) and the (T)

SOMs in the first layer into a further 4x4pixel output. The 3rd layer is composed by a single SOM merging the outputs of the 2nd layers (a 8x4pixels retina) into a 4x4pixels output. The 4rd layer gets all the outputs from the previous layers as a $(3*400+2*32+16)$ pixels flattened input vector and merges it into a 3x3pixels output. **The 4th layer gives abstracted representations of the 3-modal sensory information as 9 identifiable states.**

2.2 Goal selection

The goal-selection layer is composed by a matrix of 3x3 units (same numerosity of the goal-abstraction layer). At the start of each trial one of the units is activated. The choice about which unit has to be activated is stochastic. The probability for each unit is given by applying a softmax to the amplitudes of 3x3 running averages of the prediction error (see section 2.3). Each average is updated at the end of trials in which the corresponding goal-selection unit is chosen. As a result these running averages convey information about the current amount of error related to each goal-selection unit. **Thus the probability of a goal to be selected is related to the amount of prediction error on it (intrinsic motivation to learn what has not yet been learned).**

2.3 Goal prediction

The **goal-prediction layer** takes the activation of the goal-selection layer as an input and outputs a **prediction of a *match* event**. This prediction is compared with the actual occurrence of a *match* event at the end of each trial. The result of this comparison (*error*) is used both to update the parameters of the goal-prediction layer and to decide the activation of the goal-selection layer (see section 2.2).

2.4 Action triggering

The **motor output** of the controller is represented by the activation of 6 readout units of an echo-state network (ESN). Each readout units gives the current amplitude of a joint angle of the 2-arm agent.




Figure 3: The echo-state network. A) General schema of the functioning of an ESN. Inputs reach sparsely the internal units of the reservoir. Internal units are connected with each other via sparse random connections. Learning consists in the update of the external weights connecting the reservoir to one or more readout units. After learning the readout units reproduce a learned trajectory in relation to a specific input. B) The typical activation of the reservoir in the model. During the second half of the trial the ESN reaches a fixed point that is peculiar of for the current input (a selected goal).

The input to the ESN is given by the goal-selection layer. The activation of the goal-selection layer is connected to the ESN in two ways:

1. It steadily excites distinct subpopulations of the ESN reservoir based on the current selection during a trial. This excitation triggers the proper dynamics of the reservoir and guarantees that the reservoir activity fades to a distinct fixed point depending on the selection.
2. It inhibits distinct subpopulations of the ESN reservoir based on the current selection. This inhibition allows to select different dynamics in the network.

The weights connecting the reservoir to the readout units of the ESN are updated via a reward-based online learning rule. **The *match* event triggers the reward signal.**

The actual motor output is composed of the activation of the 6 readout units and that of 6 sinusoidal oscillators whose parameters are set randomly at each trial. These oscillators add **exploratory** noise to the motor output. Their amplitude depends on the amount of previous *match* events given the same selection. As soon as matches become frequent the motor output becomes completely dependent on the activity of the readout units (**exploitation**).

3 Simulation flow

Each trial is composed of two phases:

- **reset interval** no goal is selected, no update is made to the weights of the motor readout units and no goal input is given to the ESN.
- **activation** all components are switched on. A goal is chosen in the goal-selection layer and is maintained fixed throughout the trial.

A trial ends if a determined amount of time is reached or if a match event occurs. At the end of each trial the prediction error is computed. The moving average of the prediction error for the current selection is updated.